Overlap encodings

Hervé Pagès

Last modified: September 2013; Compiled: October 15, 2013

Contents

1   Introduction

In the context of an RNA-seq experiment, encoding the overlaps between the aligned reads and the transcripts can be used for detecting those overlaps that are "compatible" with the splicing of the transcript.

Various tools are provided in the *IRanges* and *GenomicRanges* packages for working with *overlap encodings*. In this vignette, we illustrate the use of these tools on the single-end and paired-end reads of an RNA-seq experiment.

2   Load reads from a BAM file

2.1   Load single-end reads from a BAM file

BAM file `untreated1_chr4.bam` (located in the *pasillaBamSubset* data package) contains single-end reads from the "Pasilla" experiment and aligned against the dm3 genome (see `?untreated1_chr4` in the *pasillaBamSubset* package for more information about those reads):

```
> library(pasillaBamSubset)
> untreated1_chr4()
```

```
[1] "/home/biocbuild/bbs-2.13-bioc/R/library/pasillaBamSubset/extdata/untreated1_chr4.bam"
```

We use the `readGAlignments` function defined in the *GenomicRanges* package to load the reads into a *GAlignments* object. It's probably a good idea to get rid of the PCR or optical duplicates (flag bit 0x400 in the SAM format, see the SAM Spec [1] for the details), as well as reads not passing quality controls (flag bit 0x200 in the SAM format). We do this by creating a *ScanBamParam* object that we pass to `readGAlignments` (see `?ScanBamParam` in the *Rsamtools* package for the details). Note that we also use `use.names=TRUE` in order to load the *query names* (aka *query template names*, see QNAME field in the SAM Spec) from the BAM file (`readGAlignments` will use them to set the names of the returned object):

```
> library(GenomicRanges)
> library(Rsamtools)
> flag0 <- scanBamFlag(isDuplicate=FALSE, isNotPassingQualityControls=FALSE)
> param0 <- ScanBamParam(flag=flag0)
```

---

[1] http://samtools.sourceforge.net/

```
> U1.GAL <- readGAlignments(untreated1_chr4(), use.names=TRUE, param=param0)
> head(U1.GAL)
```

```
GAlignments with 6 alignments and 0 metadata columns:
                    seqnames strand      cigar    qwidth     start       end     width      ngap
                       <Rle>  <Rle> <character> <integer> <integer> <integer> <integer> <integer>
  SRR031729.3941844     chr4      -         75M        75       892       966        75         0
  SRR031728.3674563     chr4      -         75M        75       919       993        75         0
  SRR031729.8532600     chr4      +         75M        75       924       998        75         0
  SRR031729.2779333     chr4      +         75M        75       936      1010        75         0
  SRR031728.2826481     chr4      +         75M        75       949      1023        75         0
  SRR031728.2919098     chr4      -         75M        75       967      1041        75         0
  ---
  seqlengths:
      chr2L     chr2R     chr3L     chr3R      chr4      chrM      chrX   chrYHet
   23011544  21146708  24543557  27905053   1351857     19517  22422827    347038
```

Because the aligner used to align those reads can report more than 1 alignment per *original query* (i.e. per read stored in the input file, typically a FASTQ file), we shouldn't expect the names of U1.GAL to be unique:

```
> U1.GAL_names_is_dup <- duplicated(names(U1.GAL))
> table(U1.GAL_names_is_dup)
```

```
U1.GAL_names_is_dup
 FALSE    TRUE
190770   13585
```

Storing the *query names* in a factor will be useful as we will see later in this document:

```
> U1.uqnames <- unique(names(U1.GAL))
> U1.GAL_qnames <- factor(names(U1.GAL), levels=U1.uqnames)
```

Note that we explicitly provide the levels of the factor to enforce their order. Otherwise factor() would put them in lexicographic order which is not advisable because it depends on the locale in use.

Another object that will be useful to keep near at hand is the mapping between each *query name* and its first occurence in U1.GAL_qnames:

```
> U1.GAL_dup2unq <- match(U1.GAL_qnames, U1.GAL_qnames)
```

Our reads can have up to 2 gaps (a gap corresponds to an N operation in the CIGAR):

```
> head(unique(cigar(U1.GAL)))
```

```
[1] "75M"          "35M6727N40M"  "22M6727N53M"  "13M6727N62M"  "26M292N49M"   "62M21227N13M"
```

```
> table(ngap(U1.GAL))
```

```
     0      1      2
184039  20169    147
```

Also, the following table indicates that indels were not allowed/supported during the alignment process (no I or D CIGAR operations):

```
> colSums(cigarOpTable(cigar(U1.GAL)))
```

```
        M        I        D        N        S        H        P        =        X
 15326625        0        0 21682582        0        0        0        0        0
```

2.2   Load paired-end reads from a BAM file

BAM file untreated3_chr4.bam (located in the *pasillaBamSubset* data package) contains paired-end reads from the
"Pasilla" experiment and aligned against the dm3 genome (see ?untreated3_chr4 in the *pasillaBamSubset* package for
more information about those reads). We use the readGAlignmentPairs function to load them into a *GAlignmentPairs*
object:

```
> U3.galp <- readGAlignmentPairs(untreated3_chr4(), use.names=TRUE, param=param0)
> head(U3.galp)

GAlignmentPairs with 6 alignment pairs and 0 metadata columns:
                  seqnames strand :       ranges --        ranges
                     <Rle>  <Rle> :    <IRanges> --     <IRanges>
SRR031715.1138209     chr4      + : [169,  205] -- [ 326,  362]
 SRR031714.756385     chr4      + : [943,  979] -- [1086, 1122]
SRR031714.2355189     chr4      + : [944,  980] -- [1119, 1155]
SRR031714.5054563     chr4      + : [946,  982] -- [ 986, 1022]
SRR031715.1722593     chr4      + : [966, 1002] -- [1108, 1144]
SRR031715.2202469     chr4      + : [966, 1002] -- [1114, 1150]
---
seqlengths:
     chr2L     chr2R     chr3L     chr3R      chr4      chrM      chrX   chrYHet
  23011544  21146708  24543557  27905053   1351857     19517  22422827    347038
```

The show method for *GAlignmentPairs* objects displays two ranges columns, one for the *first* alignment in the pair
(the left column), and one for the *last* alignment in the pair (the right column). The strand column corresponds to the
strand of the *first* alignment.

```
> head(first(U3.galp))

GAlignments with 6 alignments and 0 metadata columns:
                  seqnames strand       cigar     qwidth       start         end       width        ngap
                     <Rle>  <Rle> <character> <integer> <integer> <integer> <integer> <integer>
SRR031715.1138209     chr4      +         37M        37         169         205          37           0
 SRR031714.756385     chr4      +         37M        37         943         979          37           0
SRR031714.2355189     chr4      +         37M        37         944         980          37           0
SRR031714.5054563     chr4      +         37M        37         946         982          37           0
SRR031715.1722593     chr4      +         37M        37         966        1002          37           0
SRR031715.2202469     chr4      +         37M        37         966        1002          37           0
---
seqlengths:
     chr2L     chr2R     chr3L     chr3R      chr4      chrM      chrX   chrYHet
  23011544  21146708  24543557  27905053   1351857     19517  22422827    347038

> head(last(U3.galp))

GAlignments with 6 alignments and 0 metadata columns:
                  seqnames strand       cigar     qwidth       start         end       width        ngap
                     <Rle>  <Rle> <character> <integer> <integer> <integer> <integer> <integer>
SRR031715.1138209     chr4      -         37M        37         326         362          37           0
 SRR031714.756385     chr4      -         37M        37        1086        1122          37           0
SRR031714.2355189     chr4      -         37M        37        1119        1155          37           0
SRR031714.5054563     chr4      -         37M        37         986        1022          37           0
SRR031715.1722593     chr4      -         37M        37        1108        1144          37           0
SRR031715.2202469     chr4      -         37M        37        1114        1150          37           0
---
```

```
  seqlengths:
      chr2L     chr2R     chr3L     chr3R      chr4      chrM      chrX   chrYHet
   23011544  21146708  24543557  27905053   1351857     19517  22422827    347038
```

According to the SAM format specifications, the aligner is expected to mark each alignment pair as *proper* or not (flag bit 0x2 in the SAM format). The SAM Spec only says that a pair is *proper* if the *first* and *last* alignments in the pair are "properly aligned according to the aligner". So the exact criteria used for setting this flag is left to the aligner.

We use isProperPair to extract this flag from the *GAlignmentPairs* object:

```
> table(isProperPair(U3.galp))
```

```
FALSE   TRUE
29518  45828
```

Even though we could do *overlap encodings* with the full object, we keep only the *proper* pairs for our downstream analysis:

```
> U3.GALP <- U3.galp[isProperPair(U3.galp)]
```

Because the aligner used to align those reads can report more than 1 alignment per *original query template* (i.e. per pair of sequences stored in the input files, typically 1 FASTQ file for the *first* ends and 1 FASTQ file for the *last* ends), we shouldn't expect the names of U3.GALP to be unique:

```
> U3.GALP_names_is_dup <- duplicated(names(U3.GALP))
> table(U3.GALP_names_is_dup)
```

```
U3.GALP_names_is_dup
FALSE   TRUE
43659   2169
```

Storing the *query template names* in a factor will be useful:

```
> U3.uqnames <- unique(names(U3.GALP))
> U3.GALP_qnames <- factor(names(U3.GALP), levels=U3.uqnames)
```

as well as having the mapping between each *query template name* and its first occurence in U3.GALP_qnames:

```
> U3.GALP_dup2unq <- match(U3.GALP_qnames, U3.GALP_qnames)
```

Our reads can have up to 1 gap per end:

```
> head(unique(cigar(first(U3.GALP))))
```

```
[1] "37M"        "6M58N31M"  "25M56N12M" "19M62N18M" "29M222N8M" "9M222N28M"
```

```
> head(unique(cigar(last(U3.GALP))))
```

```
[1] "37M"        "12M58N25M"  "19M58N18M"   "27M2339N10M" "29M2339N8M"  "9M222N28M"
```

```
> table(ngap(first(U3.GALP)), ngap(last(U3.GALP)))
```

```
        0     1
  0 44510   596
  1   637    85
```

Like for our single-end reads, the following tables indicate that indels were not allowed/supported during the alignment process:

```
> colSums(cigarOpTable(cigar(first(U3.GALP))))
```

M	I	D	N	S	H	P	=	X
1695636	0	0	673919	0	0	0	0	0

```
> colSums(cigarOpTable(cigar(last(U3.GALP))))
```

M	I	D	N	S	H	P	=	X
1695636	0	0	630395	0	0	0	0	0

3   Validate the alignments produced by the aligner

In this section we show how to validate the alignments produced by the aligner by comparing the *original query sequences* (aka "true" or "real" query sequences, or query sequences **before** alignment) with the *reference query sequences* (i.e. the query sequences **after** alignment).

Note that even though this step is not strictly required for computing the *overlap encodings*, some of the concepts and string-based computations described in this section are slightly related to some ideas introduced later in this document.

3.1   Validate the single-end alignments

3.1.1   Load the *original query sequences*

To load the *original query sequences*, we reload the BAM file but now we explicitly request the SEQ field by using `what="seq"` in our call to `ScanBamParam`. To further validate the alignments produced by the aligner, we also need to load the NM tag which is a predefined tag described in the SAM Spec as the "Edit distance to the reference, including ambiguous bases but excluding clipping" (note that tags are optional fields i.e. not all BAM files have them):

```
> param1 <- ScanBamParam(flag=flag0, what="seq", tag="NM")
> U1.GAL <- readGAlignments(untreated1_chr4(), use.names=TRUE, param=param1)
> U1.GAL_qseq <- mcols(U1.GAL)$seq
> names(U1.GAL_qseq) <- names(U1.GAL)
> head(U1.GAL_qseq)

  A DNAStringSet instance of length 6
    width seq                                                         names
[1]    75 CTGTGGTGACCAACACCACAGAATGGTTCGGGC...GGGTTCCCTGCCCCTTTCCTGGCTAGGTTGTCC SRR031729.3941844
[2]    75 TCGGGCCCAATTAGAGGGTTCCCTGCCCCTTTC...TTGTCCGCTAGCTCATTTCCTGGGCTGTTGTTG SRR031728.3674563
[3]    75 CCCAATTAGAGGATTCTCTGCCCCTTTCCTGGC...CGGTAGCTCATTTCCCGGGATGTTGTTGTGTCC SRR031729.8532600
[4]    75 GTTCTCTGCCCCTTTCCTGGCTAGGTTGTCCGC...TCCCGAGATGTTGTTGTGTCCCGGGACCCACCT SRR031729.2779333
[5]    75 TTCCTGGCTAGGTTGTCCGCTAGCTCATTTCCC...TTGTGTCCCGGGACACACCTTATTGTGAGTTTG SRR031728.2826481
[6]    75 GCTAGCTCATTTCCCGGGAGGTTGTTGTGTCCC...CTTATTGTGAGTTTGTTGACAGCTCCAAGTTTG SRR031728.2919098
```

Because the BAM format imposes that the read sequence is "reverse complemented" when the read is aligned to the minus strand, we "reverse complement" it again to get the *original query sequences*:

```
> U1.GAL_oqseq <- U1.GAL_qseq
> U1.GAL_is_on_minus <- as.logical(strand(U1.GAL) == "-")
> U1.GAL_oqseq[U1.GAL_is_on_minus] <- reverseComplement(U1.GAL_oqseq[U1.GAL_is_on_minus])
> head(U1.GAL_oqseq)

  A DNAStringSet instance of length 6
    width seq                                                         names
[1]    75 GGACAACCTAGCCAGGAAAGGGGCAGGGAACCC...GCCCGAACCATTCTGTGGTGTTGGTCACCACAG SRR031729.3941844
[2]    75 CAACAACAGCCCAGGAAATGAGCTAGCGGACAA...GAAAGGGGCAGGGAACCCTCTAATTGGGCCCGA SRR031728.3674563
[3]    75 CCCAATTAGAGGATTCTCTGCCCCTTTCCTGGC...CGGTAGCTCATTTCCCGGGATGTTGTTGTGTCC SRR031729.8532600
[4]    75 GTTCTCTGCCCCTTTCCTGGCTAGGTTGTCCGC...TCCCGAGATGTTGTTGTGTCCCGGGACCCACCT SRR031729.2779333
[5]    75 TTCCTGGCTAGGTTGTCCGCTAGCTCATTTCCC...TTGTGTCCCGGGACACACCTTATTGTGAGTTTG SRR031728.2826481
[6]    75 CAAACTTGGAGCTGTCAACAAACTCACAATAAG...GGGACACAACAACCTCCCGGGAAATGAGCTAGC SRR031728.2919098
```

Note that sequences with the same *query name* correspond to the same *original query* and therefore must be the same. Let's do a quick sanity check:

```
> stopifnot(all(U1.GAL_oqseq == U1.GAL_oqseq[U1.GAL_dup2unq]))
```

Finally, let's reduce U1.GAL_oqseq to one *original query sequence* per unique *query name*:

```
> U1.oqseq <- U1.GAL_oqseq[!U1.GAL_names_is_dup]
```

If we had access to the FASTQ file used as input to the aligner, U1.oqseq would be the subset of this file made of the sequences with at least 1 alignment reported in BAM file `untreated1_chr4.bam`.

3.1.2 Compute the *reference query sequences*

The *reference query sequences* can easily be computed by extracting the nucleotides mapped to each read from the reference genome. This of course requires that we have access to the reference genome used by the aligner. In Bioconductor, the full genome sequence for the dm3 assembly is stored in the *BSgenome.Dmelanogaster.UCSC.dm3* data package [2]:

```
> library(BSgenome.Dmelanogaster.UCSC.dm3)
> Dmelanogaster

Fly genome
|
| organism: Drosophila melanogaster (Fly)
| provider: UCSC
| provider version: dm3
| release date: Apr. 2006
| release name: BDGP Release 5
|
| single sequences (see '?seqnames'):
|   chr2L      chr2R      chr3L      chr3R      chr4       chrX       chrU       chrM
|   chr2LHet   chr2RHet   chr3LHet   chr3RHet   chrXHet    chrYHet    chrUextra
|
| multiple sequences (see '?mseqnames'):
|   upstream1000  upstream2000  upstream5000
|
| (use the '$' or '[[' operator to access a given sequence)
```

Let's start by converting our *GAlignments* object U1.GAL into a *GRangesList* object:

```
> U1.grl <- grglist(U1.GAL, order.as.in.query=TRUE)
```

To extract the portions of the reference genome corresponding to the ranges in U1.grl, we can use the extractTranscriptsFromGenome function defined in the *GenomicFeatures* package:

```
> library(GenomicFeatures)
> U1.GAL_rqseq <- extractTranscriptsFromGenome(Dmelanogaster, U1.grl)
> head(U1.GAL_rqseq)

  A DNAStringSet instance of length 6
    width seq                                                              names
[1]    75 GGACAACCTAGCCAGGAAAGGGGCAGAGAACCC...GCCCGAACCATTCTGTGGTGTTGGTCACCACAG SRR031729.3941844
[2]    75 CAACAACATCCCGGGAAATGAGCTAGCGGACAA...GAAAGGGGCAGAGAACCCTCTAATTGGGCCCGA SRR031728.3674563
[3]    75 CCCAATTAGAGGGTTCTCTGCCCCTTTCCTGGC...CGCTAGCTCATTTCCCGGGATGTTGTTGTGTCC SRR031729.8532600
[4]    75 GTTCTCTGCCCCTTTCCTGGCTAGGTTGTCCGC...TCCCGGGATGTTGTTGTGTCCCGGGACCCACCT SRR031729.2779333
[5]    75 TTCCTGGCTAGGTTGTCCGCTAGCTCATTTCCC...TTGTGTCCCGGGACCCACCTTATTGTGAGTTTG SRR031728.2826481
[6]    75 CAAACTTGGAGCTGTCAACAAACTCACAATAAG...GGGACACAACAACATCCCGGGAAATGAGCTAGC SRR031728.2919098
```

3.1.3 Compare the *original query sequences* with the *reference query sequences*

We can use the neditAt function defined in the *Biostrings* package to compute the edit distance between 2 strings. Because the aligned reads have no indels, we should only see mismatches (typically a small number) during that comparison so we don't need to call neditAt with with.indels=TRUE. And because calling neditAt in a loop is inefficient, we only do this comparison for the first 500 sequences in U1.GAL_oqseq:

```
> U1.GAL_nedit500 <- sapply(1:500, function(i) neditAt(U1.GAL_oqseq[[i]], U1.GAL_rqseq[[i]]))
> table(U1.GAL_nedit500)
```

---

[2]See http://bioconductor.org/packages/release/data/annotation/ for the full list of annotation packages available in the current release of Bioconductor.

```
U1.GAL_nedit500
  0   1   2   3   4   5
273 101  97  22   6   1
```

   Yes, the first 500 sequences in `U1.GAL_oqseq` are "close" to the first 500 sequences in `U1.GAL_rqseq`.
   Now let's compare the edit distance reported by `neditAt` with the edit distance reported by the aligner (NM tag). Because the latter excludes the N CIGAR operations, it should actually be the same as the former. We confirm this for the 500 edit distances computed in `U1.GAL_nedit500`:

```
> U1.GAL_NM <- mcols(U1.GAL)$NM
> stopifnot(all(U1.GAL_NM[1:500] == U1.GAL_nedit500))
```

   Note that the maximum observed number of mismatches tells us how many mismatches per read were tolerated by the aligner:

```
> table(U1.GAL_NM)
```

```
U1.GAL_NM
     0      1      2      3      4      5      6
135911  44033  15684   5172   2233    948    374
```

3.2   Validate the paired-end alignments

3.2.1   Load the *original query sequences*

To load the *original query sequences*, we reload the BAM file and request the SEQ field (and also the NM tag). Since we've removed the improper pairs from our current U3.GALP object, we need to do this again but now we do it at load time which is equivalent to doing it afterward (i.e. not only do we have the guarantee to end up with the same elements in U3.GALP, but also to have them in the same order):

```
> flag2 <- scanBamFlag(isDuplicate=FALSE,
+                      isNotPassingQualityControls=FALSE,
+                      isProperPair=TRUE)
> param2 <- ScanBamParam(flag=flag2, what="seq", tag="NM")
> U3.GALP <- readGAlignmentPairs(untreated3_chr4(), use.names=TRUE, param=param2)
```

   Let's extract the *first* and *last* sequences from U3.GALP:

```
> U3.GALP_qseq1 <- mcols(first(U3.GALP))$seq
> U3.GALP_qseq2 <- mcols(last(U3.GALP))$seq
> names(U3.GALP_qseq1) <- names(U3.GALP_qseq2) <- names(U3.GALP)
> head(U3.GALP_qseq1)
```

```
  A DNAStringSet instance of length 6
    width seq                                                                   names
[1]    37 CCGTTTCTGAAGGAGATGGCTCATGGAGTACCTGCCT                              SRR031715.1138209
[2]    37 GCCCCTTTCCTGGCTAGGTTGTCCGCTAGCTCATTTC                              SRR031714.756385
[3]    37 CCTTTCCTGGCTAGGTTGTCCGCTAGCTCATTTCCCG                              SRR031714.5054563
[4]    37 CGCTAGCTCATTTCCCGGGCTGTTGTTGTGTCCCGGG                              SRR031715.1722593
[5]    37 CGCTAGCTCATTTCCCGAGATGTTGTTGTGTCCCGGG                              SRR031715.2202469
[6]    37 GCTTTGCTGAGCGCCTTTATGGCTGCTTGACTATCAG                              SRR031714.3544437
```

```
> head(U3.GALP_qseq2)
```

```
  A DNAStringSet instance of length 6
    width seq                                                                   names
[1]    37 GTCTCCAGCAGAGCAGATGGAGCAACGGCCTATAGAG                              SRR031715.1138209
```

```
[2]    37 AGCTTTGCTGAGCGCCTTTATGGCTGCTTGACTATCA                              SRR031714.756385
[3]    37 TGTTGTTGTGTCCCGGGACCCACCTTATTGTGAGTTT                              SRR031714.5054563
[4]    37 GCTGCTTGACTATCAGACAGTATAGCAATGTCCTTGC                              SRR031715.1722593
[5]    37 TGACTATCAGACAGTATAGCAATGTCCTTGCCATGAT                              SRR031715.2202469
[6]    37 GTCCGCTAGCTCATTTCCCGGGATGTTTTGTGTCCC                               SRR031714.3544437
```

To obtain the *original query sequences* we "reverse complement" the sequences that are aligned to the minus strand:

```
> U3.GALP_oqseq1 <- U3.GALP_qseq1
> U3.GALP_first_is_on_minus <- as.logical(strand(first(U3.GALP)) == "-")
> U3.GALP_oqseq1[U3.GALP_first_is_on_minus] <- reverseComplement(U3.GALP_oqseq1[U3.GALP_first_is_on_minus],
> U3.GALP_oqseq2 <- U3.GALP_qseq2
> U3.GALP_last_is_on_minus <- as.logical(strand(last(U3.GALP)) == "-")
> U3.GALP_oqseq2[U3.GALP_last_is_on_minus] <- reverseComplement(U3.GALP_oqseq2[U3.GALP_last_is_on_minus])
```

Note that sequence pairs with the same *query template name* correspond to the same *original query pairs* and therefore should be the same. Let's do a quick sanity check:

```
> stopifnot(all(U3.GALP_oqseq1 == U3.GALP_oqseq1[U3.GALP_dup2unq]))
> stopifnot(all(U3.GALP_oqseq2 == U3.GALP_oqseq2[U3.GALP_dup2unq]))
```

Finally, let's reduce U3.GALP_oqseq1 and U3.GALP_oqseq2 to one *original query sequence* per unique *query template name*:

```
> U3.oqseq1 <- U3.GALP_oqseq1[!U3.GALP_names_is_dup]
> U3.oqseq2 <- U3.GALP_oqseq2[!U3.GALP_names_is_dup]
```

If we had access to the 2 FASTQ files used as input to the aligner, U3.oqseq1 and U3.oqseq2 would be the subsets of those files made of the sequence pairs with at least 1 alignment pair reported in BAM file untreated3_chr4.bam.

3.2.2 Compute the *reference query sequences*

Because our reads are paired-end, we extract separately the ranges corresponding to their *first* ends (aka *first* segments in BAM jargon) and those corresponding to their *last* ends (aka *last* segments in BAM jargon):

```
> U3.grl_first <- grglist(first(U3.GALP), order.as.in.query=TRUE)
> U3.grl_last <- grglist(last(U3.GALP, invert.strand=TRUE), order.as.in.query=TRUE)
```

Then we extract the portions of the reference genome corresponding to the ranges in *GRangesList* objects U3.grl_first and U3.grl_last:

```
> U3.GALP_rqseq1 <- extractTranscriptsFromGenome(Dmelanogaster, U3.grl_first)
> U3.GALP_rqseq2 <- extractTranscriptsFromGenome(Dmelanogaster, U3.grl_last)
```

3.2.3 Compare the *original query sequences* with the *reference query sequences*

Because the aligned reads have no indels, we should only see mismatches (typically a small number) during that comparison so we don't need to call neditAt with with.indels=TRUE. Let's do this comparison for the first 500 sequences in U3.GALP_oqseq1 and in reverseComplement(U3.GALP_oqseq2):

```
> U3.GALP_first_nedit500 <- sapply(1:500, function(i)
+     neditAt(U3.GALP_oqseq1[[i]], U3.GALP_rqseq1[[i]])
+ )
> table(U3.GALP_first_nedit500)

U3.GALP_first_nedit500
  0   1   2
337 135  28
```

```
> U3.GALP_last_nedit500 <- sapply(1:500, function(i)
+     neditAt(reverseComplement(U3.GALP_oqseq2[[i]]), U3.GALP_rqseq2[[i]])
+ )
> table(U3.GALP_last_nedit500)

U3.GALP_last_nedit500
  0   1   2
341 121  38
```

Yes, the first 500 sequences in U3.GALP_oqseq1 and in reverseComplement(U3.GALP_oqseq2) are "close" to the first 500 sequences in U3.GALP_rqseq1 and in U3.GALP_rqseq2, respectively.

Now let's compare the edit distance reported by `neditAt` with the edit distance reported by the aligner (NM tag). Because the latter excludes the N CIGAR operations, it should actually be the same as the former. We confirm this for the 500 edit distances computed in U3.GALP_first_nedit500 and U3.GALP_last_nedit500:

```
> U3.GALP_first_NM <- mcols(first(U3.GALP))$NM
> stopifnot(all(U3.GALP_first_NM[1:500] == U3.GALP_first_nedit500))
> U3.GALP_last_NM <- mcols(last(U3.GALP))$NM
> stopifnot(all(U3.GALP_last_NM[1:500] == U3.GALP_last_nedit500))
```

Note that the following table tells us how many mismatches per read were tolerated by the aligner:

```
> table(U3.GALP_first_NM, U3.GALP_last_NM)

                U3.GALP_last_NM
U3.GALP_first_NM     0     1     2
               0 30140  4648  1292
               1  5023  1522   760
               2  1240   700   503
```

Up to 2 mismatches per end.

3.3   Conclusion

In addition to validate the alignments produced by the aligner, the validation described in this section is also an efficient and accurate way to make sure that the reference genome we've picked up is the same as the reference genome used by the aligner, at least for the regions covered by the reads.

In other words, if it's known that the 2 reference genomes are different, then this validation could still be performed, and, if successful, would indicate that the 2 genomes are probably substitutable for most analysis happening downstream of the BAM file.

4   Find all the overlaps between the reads and transcripts

4.1   Load the transcripts from a TranscriptDb object

In order to compute overlaps between reads and transcripts, we need access to the genomic positions of a set of known transcripts and their exons. It is essential that the reference genome of this set of transcripts and exons be **exactly** the same as the reference genome used to align the reads.

We could use the `makeTranscriptDbFromUCSC` function defined in the *GenomicFeatures* package to make a *TranscriptDb* object containing the dm3 transcripts and their exons retrieved from the UCSC Genome Browser[3]. The Bioconductor project however provides a few annotation packages containing *TranscriptDb* objects for the most commonly studied organisms (those data packages are sometimes called the *TxDb* packages). One of them is the *TxDb.Dmelanogaster.-UCSC.dm3.ensGene* package. It contains a *TranscriptDb* object that was made by pointing the `makeTranscriptDbFromUCSC` function to the dm3 genome and *Ensembl Genes* track [4]. We can use it here:

---

[3]http://genome.ucsc.edu/cgi-bin/hgGateway

[4]See http://genome.ucsc.edu/cgi-bin/hgTrackUi?hgsid=276880911&g=ensGene for a description of this track.

```
> library(TxDb.Dmelanogaster.UCSC.dm3.ensGene)
> TxDb.Dmelanogaster.UCSC.dm3.ensGene

TranscriptDb object:
| Db type: TranscriptDb
| Supporting package: GenomicFeatures
| Data source: UCSC
| Genome: dm3
| Organism: Drosophila melanogaster
| UCSC Table: ensGene
| Resource URL: http://genome.ucsc.edu/
| Type of Gene ID: Ensembl gene ID
| Full dataset: yes
| miRBase build ID: NA
| transcript_nrow: 29173
| exon_nrow: 76920
| cds_nrow: 62135
| Db created by: GenomicFeatures package from Bioconductor
| Creation time: 2013-09-19 00:24:08 -0700 (Thu, 19 Sep 2013)
| GenomicFeatures version at creation time: 1.13.40
| RSQLite version at creation time: 0.11.4
| DBSCHEMAVERSION: 1.0

> txdb <- TxDb.Dmelanogaster.UCSC.dm3.ensGene
```

We extract the exons grouped by transcript in a *GRangesList* object:

```
> exbytx <- exonsBy(txdb, by="tx", use.names=TRUE)
> length(exbytx)  # nb of transcripts

[1] 29173
```

We check that all the exons in any given transcript belong to the same chromosome and strand. Knowing that our set of transcripts is free of this sort of trans-splicing events typically allows some significant simplifications during the downstream analysis [5]. A quick and easy way to check this is to take advantage of the fact that `seqnames` and `strand` return *RleList* objects. So we can extract the number of Rle runs for each transcript and make sure it's always 1:

```
> table(elementLengths(runLength(seqnames(exbytx))))

    1
29173

> table(elementLengths(runLength(strand(exbytx))))

    1
29173
```

Therefore the strand of any given transcript is unambiguously defined and can be extracted with:

```
> exbytx_strand <- unlist(runValue(strand(exbytx)), use.names=FALSE)
```

We will also need the mapping between the transcripts and their gene. We start by using `transcripts` to extract this information from our *TranscriptDb* object txdb, and then we construct a named factor that represents the mapping:

```
> tx <- transcripts(txdb, columns=c("tx_name", "gene_id"))
> head(tx)
```

---
[5]Dealing with trans-splicing events is not covered in this document.

```
GRanges with 6 ranges and 2 metadata columns:
      seqnames           ranges strand |    tx_name         gene_id
         <Rle>        <IRanges>  <Rle> | <character> <CharacterList>
  [1]    chr2L [ 7529,  9484]      + | FBtr0300689     FBgn0031208
  [2]    chr2L [ 7529,  9484]      + | FBtr0300690     FBgn0031208
  [3]    chr2L [ 7529,  9484]      + | FBtr0330654     FBgn0031208
  [4]    chr2L [21952, 24237]      + | FBtr0309810     FBgn0263584
  [5]    chr2L [66584, 71390]      + | FBtr0306539     FBgn0067779
  [6]    chr2L [67043, 71081]      + | FBtr0306536     FBgn0067779
  ---
  seqlengths:
      chr2L     chr2R     chr3L     chr3R      chr4 ...   chr3RHet    chrXHet   chrYHet chrUextra
   23011544  21146708  24543557  27905053   1351857 ...    2517507     204112    347038  29004656
```

```
> df <- mcols(tx)
> exbytx2gene <- as.character(df$gene_id)
> exbytx2gene <- factor(exbytx2gene, levels=unique(exbytx2gene))
> names(exbytx2gene) <- df$tx_name
> exbytx2gene <- exbytx2gene[names(exbytx)]
> head(exbytx2gene)
```

```
FBtr0300689 FBtr0300690 FBtr0330654 FBtr0309810 FBtr0306539 FBtr0306536
FBgn0031208 FBgn0031208 FBgn0031208 FBgn0263584 FBgn0067779 FBgn0067779
15682 Levels: FBgn0031208 FBgn0263584 FBgn0067779 FBgn0031213 FBgn0031214 FBgn0031216 ... FBgn0085792
```

```
> nlevels(exbytx2gene)  # nb of genes
```

```
[1] 15682
```

4.2   Single-end overlaps

4.2.1   Find the single-end overlaps

We are ready to compute the overlaps with the `findOverlaps` function. Note that the strand of the queries produced by the RNA-seq experiment is typically unknown so we use `ignore.strand=TRUE`:

```
> U1.OV00 <- findOverlaps(U1.GAL, exbytx, ignore.strand=TRUE)
```

U1.OV00 is a *Hits* object that contains 1 element per overlap. Its length gives the number of overlaps:

```
> length(U1.OV00)
```

```
[1] 563552
```

4.2.2   Tabulate the single-end overlaps

We will repeatedly use the 2 following little helper functions to "tabulate" the overlaps in a given *Hits* object (e.g. U1.OV00), i.e. to count the number of overlaps for each element in the query or for each element in the subject:

```
> nhitPerQuery <- function(x) tabulate(queryHits(x), nbins=queryLength(x))
> nhitPerSubject <- function(x) tabulate(subjectHits(x), nbins=subjectLength(x))
```

Number of transcripts for each alignment in U1.GAL:

```
> U1.GAL_ntx <- nhitPerQuery(U1.OV00)
> mcols(U1.GAL)$ntx <- U1.GAL_ntx
> head(U1.GAL)
```

```
GAlignments with 6 alignments and 3 metadata columns:
                    seqnames strand       cigar    qwidth     start       end     width      ngap |
                       <Rle>  <Rle> <character> <integer> <integer> <integer> <integer> <integer> |
  SRR031729.3941844      chr4      -         75M        75       892       966        75         0 |
  SRR031728.3674563      chr4      -         75M        75       919       993        75         0 |
  SRR031729.8532600      chr4      +         75M        75       924       998        75         0 |
  SRR031729.2779333      chr4      +         75M        75       936      1010        75         0 |
  SRR031728.2826481      chr4      +         75M        75       949      1023        75         0 |
  SRR031728.2919098      chr4      -         75M        75       967      1041        75         0 |
                                                                                              seq
                                                                                    <DNAStringSet>
  SRR031729.3941844 CTGTGGTGACCAACACCACAGAATGGTTCGGGCCCAATTAGAGGGTTCCCTGCCCCTTTCCTGGCTAGGTTGTCC
  SRR031728.3674563 TCGGGCCCAATTAGAGGGTTCCCTGCCCCTTTCCTGGCTAGGTTGTCCGCTAGCTCATTTCCTGGGCTGTTGTTG
  SRR031729.8532600 CCCAATTAGAGGATTCTCTGCCCCTTTCCTGGCTAGGTTGTCCGGTAGCTCATTTCCCGGGATGTTGTTGTGTCC
  SRR031729.2779333 GTTCTCTGCCCCTTTCCTGGCTAGGTTGTCCGCTAGCTCATTTCCCGAGATGTTGTTGTGTCCCGGGACCCACCT
  SRR031728.2826481 TTCCTGGCTAGGTTGTCCGCTAGCTCATTTCCCGGGCTGTTGTTGTGTCCCGGGACACACCTTATTGTGAGTTTG
  SRR031728.2919098 GCTAGCTCATTTCCCGGGAGGTTGTTGTGTCCCGGGACCCACCTTATTGTGAGTTTGTTGACAGCTCCAAGTTTG
                           NM       ntx
                    <integer> <integer>
  SRR031729.3941844         1         0
  SRR031728.3674563         3         0
  SRR031729.8532600         2         0
  SRR031729.2779333         1         0
  SRR031728.2826481         2         0
  SRR031728.2919098         1         0
  ---
  seqlengths:
       chr2L     chr2R     chr3L     chr3R      chr4      chrM      chrX   chrYHet
    23011544  21146708  24543557  27905053   1351857     19517  22422827    347038

> table(U1.GAL_ntx)

U1.GAL_ntx
    0     1     2     3     4     5     6     7     8     9    10    11    12
47076  9493 26146 82427  5291 14530  8158   610  1952  2099   492  4945  1136

> mean(U1.GAL_ntx >= 1)

[1] 0.7696362
```

76% of the alignments in U1.GAL have an overlap with at least 1 transcript in exbytx.
Note that countOverlaps can be used directly on U1.GAL and exbytx for computing U1.GAL_ntx:

```
> U1.GAL_ntx_again <- countOverlaps(U1.GAL, exbytx, ignore.strand=TRUE)
> stopifnot(identical(unname(U1.GAL_ntx_again), U1.GAL_ntx))
```

Because U1.GAL can (and actually does) contain more than 1 alignment per *original query* (aka read), we also count
the number of transcripts for each read:

```
> U1.OV10 <- remapHits(U1.OV00, query.map=U1.GAL_qnames)
> U1.uqnames_ntx <- nhitPerQuery(U1.OV10)
> names(U1.uqnames_ntx) <- U1.uqnames
> table(U1.uqnames_ntx)

U1.uqnames_ntx
    0     1     2     3     4     5     6     7     8     9    10    11    12
39503  9298 18394 82346  5278 14536  9208   610  2930  2099   488  4944  1136
```

```
> mean(U1.uqnames_ntx >= 1)
```

```
[1] 0.7929287
```

> 78.4% of the reads have an overlap with at least 1 transcript in exbytx.
> Number of reads for each transcript:

```
> U1.exbytx_nOV10 <- nhitPerSubject(U1.OV10)
> names(U1.exbytx_nOV10) <- names(exbytx)
> mean(U1.exbytx_nOV10 >= 50)
```

```
[1] 0.009015185
```

> Only 0.869% of the transcripts in exbytx have an overlap with at least 50 reads.
> Top 10 transcripts:

```
> head(sort(U1.exbytx_nOV10, decreasing=TRUE), n=10)
```

```
FBtr0308296 FBtr0089175 FBtr0089176 FBtr0112904 FBtr0289951 FBtr0089243 FBtr0333672 FBtr0089186
      40654       40529       40529       11735       11661       11656       10087       10084
FBtr0089187 FBtr0089172
      10084        6749
```

4.3   Paired-end overlaps

4.3.1   Find the paired-end overlaps

Like with our single-end overlaps, we call findOverlaps with ignore.strand=TRUE:

```
> U3.OV00 <- findOverlaps(U3.GALP, exbytx, ignore.strand=TRUE)
```

> Like U1.OV00, U3.OV00 is a *Hits* object. Its length gives the number of paired-end overlaps:

```
> length(U3.OV00)
```

```
[1] 113827
```

4.3.2   Tabulate the paired-end overlaps

Number of transcripts for each alignment pair in U3.GALP:

```
> U3.GALP_ntx <- nhitPerQuery(U3.OV00)
> mcols(U3.GALP)$ntx <- U3.GALP_ntx
> head(U3.GALP)
```

```
GAlignmentPairs with 6 alignment pairs and 1 metadata column:
                    seqnames strand :        ranges --        ranges |        ntx
                       <Rle>  <Rle> :     <IRanges> --     <IRanges> |  <integer>
SRR031715.1138209       chr4      + : [ 169,  205] -- [ 326,  362] |          0
  SRR031714.756385      chr4      + : [ 943,  979] -- [1086, 1122] |          0
SRR031714.5054563       chr4      + : [ 946,  982] -- [ 986, 1022] |          0
SRR031715.1722593       chr4      + : [ 966, 1002] -- [1108, 1144] |          0
SRR031715.2202469       chr4      + : [ 966, 1002] -- [1114, 1150] |          0
SRR031714.3544437       chr4      - : [1087, 1123] -- [ 963,  999] |          0
  ---
  seqlengths:
      chr2L      chr2R      chr3L      chr3R       chr4       chrM       chrX    chrYHet
   23011544   21146708   24543557   27905053    1351857      19517   22422827     347038
```

```
> table(U3.GALP_ntx)

U3.GALP_ntx
    0     1     2     3     4     5     6     7     8     9    10    11    12
12950  2080  5854 17025  1078  3083  2021    70   338   370    59   803    97

> mean(U3.GALP_ntx >= 1)

[1] 0.7174217
```

   71% of the alignment pairs in U3.GALP have an overlap with at least 1 transcript in exbytx.
   Note that countOverlaps can be used directly on U3.GALP and exbytx for computing U3.GALP_ntx:

```
> U3.GALP_ntx_again <- countOverlaps(U3.GALP, exbytx, ignore.strand=TRUE)
> stopifnot(identical(unname(U3.GALP_ntx_again), U3.GALP_ntx))
```

   Because U3.GALP can (and actually does) contain more than 1 alignment pair per *original query template*, we also
count the number of transcripts for each template:

```
> U3.OV10 <- remapHits(U3.OV00, query.map=U3.GALP_qnames)
> U3.uqnames_ntx <- nhitPerQuery(U3.OV10)
> names(U3.uqnames_ntx) <- U3.uqnames
> table(U3.uqnames_ntx)

U3.uqnames_ntx
    0     1     2     3     4     5     6     7     8     9    10    11    12
11851  2061  4289 17025  1193  3084  2271    70   486   370    59   803    97

> mean(U3.uqnames_ntx >= 1)

[1] 0.7285554
```

   72.3% of the templates have an overlap with at least 1 transcript in exbytx.
   Number of templates for each transcript:

```
> U3.exbytx_nOV10 <- nhitPerSubject(U3.OV10)
> names(U3.exbytx_nOV10) <- names(exbytx)
> mean(U3.exbytx_nOV10 >= 50)

[1] 0.00712988
```

   Only 0.756% of the transcripts in exbytx have an overlap with at least 50 templates.
   Top 10 transcripts:

```
> head(sort(U3.exbytx_nOV10, decreasing=TRUE), n=10)

FBtr0308296 FBtr0089175 FBtr0089176 FBtr0112904 FBtr0089243 FBtr0289951 FBtr0333672 FBtr0089186
       7574        7573        7572        2750        2732        2732        2260        2260
FBtr0089187 FBtr0310542
       2260        1698
```

5   Encode the overlaps between the reads and transcripts

5.1   Single-end encodings

The *overlap encodings* are strand sensitive so we will compute them twice, once for the "original alignments" (i.e. the
alignments of the *original queries*), and once again for the "flipped alignments" (i.e. the alignments of the "flipped *original
queries*"). We extract the ranges of the "original" and "flipped" alignments in 2 *GRangesList* objects with:

```
> U1.grlf <- flipQuery(U1.grl)  # flipped
```

and encode their overlaps with the transcripts:

```
> U1.ovencA <- encodeOverlaps(U1.grl, exbytx, hits=U1.OV00)
> U1.ovencB <- encodeOverlaps(U1.grlf, exbytx, hits=U1.OV00)
```

U1.ovencA and U1.ovencB are 2 *OverlapsEncodings* objects of the same length as *Hits* object U1.OV00. For each hit in U1.OV00, we have 2 corresponding encodings, one in U1.ovencA and one in U1.ovencB, but only one of them encodes a hit between alignment ranges and exon ranges that are on the same strand. We use the selectEncoding-WithCompatibleStrand function to merge them into a single *OverlapsEncodings* of the same length. For each hit in U1.OV00, this selects the encoding corresponding to alignment ranges and exon ranges with compatible strand:

```
> U1.grl_strand <- unlist(runValue(strand(U1.grl)), use.names=FALSE)
> U1.ovenc <- selectEncodingWithCompatibleStrand(U1.ovencA, U1.ovencB,
+                                                U1.grl_strand, exbytx_strand,
+                                                hits=U1.OV00)
> U1.ovenc

OverlapEncodings of length 563552
          Loffset Roffset encoding flippedQuery
[1]             0       3    1:i:           TRUE
[2]             4       0    1:k:          FALSE
[3]             4       0    1:k:           TRUE
[4]             4       0    1:k:           TRUE
[5]             4       0    1:k:           TRUE
[6]             4       0    1:k:           TRUE
[7]             4       0    1:k:           TRUE
[8]             4       0    1:i:           TRUE
[9]             4       0    1:i:           TRUE
...           ...     ...     ...            ...
[563544]       23       0    1:i:          FALSE
[563545]       24       0    1:i:          FALSE
[563546]       24       0    1:i:          FALSE
[563547]       23       0    1:i:          FALSE
[563548]       22       0    1:i:           TRUE
[563549]       23       0    1:i:           TRUE
[563550]       24       0    1:i:           TRUE
[563551]       24       0    1:i:           TRUE
[563552]       23       0    1:i:           TRUE
```

As a convenience, the 2 above calls to encodeOverlaps + merging step can be replaced by a single call to en-codeOverlaps on U1.grl (or U1.grlf) with flip.query.if.wrong.strand=TRUE:

```
> U1.ovenc_again <- encodeOverlaps(U1.grl, exbytx, hits=U1.OV00, flip.query.if.wrong.strand=TRUE)
> stopifnot(identical(U1.ovenc_again, U1.ovenc))
```

Unique encodings in U1.ovenc:

```
> U1.unique_encodings <- levels(U1.ovenc)
> length(U1.unique_encodings)

[1] 120

> head(U1.unique_encodings)

[1] "1:c:" "1:e:" "1:f:" "1:h:" "1:i:" "1:j:"
```

```
> U1.ovenc_table <- table(encoding(U1.ovenc))
> tail(sort(U1.ovenc_table))

    1:f:    1:k:c:     1:k:     1:c: 2:jm:af:     1:i:
    1555      1889     8800     9523    72929   455176
```

Encodings are sort of cryptic but utilities are provided to extract specific meaning from them. Use of these utilities is covered later in this document.

5.2 Paired-end encodings

Let's encode the overlaps in U3.OV00:

```
> U3.grl <- grglist(U3.GALP, order.as.in.query=TRUE)
> U3.ovenc <- encodeOverlaps(U3.grl, exbytx, hits=U3.OV00, flip.query.if.wrong.strand=TRUE)
> U3.ovenc

OverlapEncodings of length 113827
          Loffset Roffset    encoding flippedQuery
[1]             4       0 1--1:i--k:         TRUE
[2]             4       0 1--1:i--i:         TRUE
[3]             4       0 1--1:i--k:        FALSE
[4]             4       0 1--1:i--k:        FALSE
[5]             4       0 1--1:a--c:         TRUE
[6]             4       0 1--1:i--m:        FALSE
[7]             3       1 1--1:i--i:        FALSE
[8]             3       1 1--1:i--i:        FALSE
[9]             2       2 1--1:i--i:         TRUE
...           ...     ...         ...          ...
[113819]       23       0 1--1:i--i:        FALSE
[113820]       24       0 1--1:i--i:        FALSE
[113821]       24       0 1--1:i--i:        FALSE
[113822]       23       0 1--1:i--i:        FALSE
[113823]       22       0 1--1:i--i:         TRUE
[113824]       23       0 1--1:i--i:         TRUE
[113825]       24       0 1--1:i--i:         TRUE
[113826]       24       0 1--1:i--i:         TRUE
[113827]       23       0 1--1:i--i:         TRUE
```

Unique encodings in U3.ovenc:

```
> U3.unique_encodings <- levels(U3.ovenc)
> length(U3.unique_encodings)

[1] 123

> head(U3.unique_encodings)

[1] "1--1:a--c:" "1--1:a--i:" "1--1:a--j:" "1--1:a--k:" "1--1:b--i:" "1--1:b--k:"

> U3.ovenc_table <- table(encoding(U3.ovenc))
> tail(sort(U3.ovenc_table))

       1--1:i--m:        1--1:i--k:        1--1:c--i: 1--2:i--jm:a--af: 2--1:jm--m:af--i:
              852              1485              1714              2480              2700
       1--1:i--i:
           100084
```

6 "Compatible" overlaps

We are interested in a particular type of overlap where the read overlaps the transcript in a "compatible" way, that is, in a way compatible with the splicing of the transcript. The isCompatibleWithSplicing function can be used on an *OverlapEncodings* object to detect this type of overlap. Note that isCompatibleWithSplicing can also be used on a character vector or factor.

6.1 "Compatible" single-end overlaps

6.1.1 "Compatible" single-end encodings

U1.ovenc contains 7 unique encodings "compatible" with the splicing of the transcript:

```
> sort(U1.ovenc_table[isCompatibleWithSplicing(U1.unique_encodings)])

     2:jm:ag:        2:gm:af: 3:jmm:agm:aaf:            1:j:            1:f:        2:jm:af:
          32              79            488            1538            1555           72929
        1:i:
      455176
```

Encodings "1:i:" (455176 occurences in U1.ovenc), "2:jm:af:" (72929 occurences in U1.ovenc), and "3:jmm:agm:aaf:" (488 occurences in U1.ovenc), correspond to the following overlaps:

- "1:i:"

    ```
    - read (no gap):            oooooooo
    - transcript:       ...   >>>>>>>>>>>>>    ...
    ```

- "2:jm:af:"

    ```
    - read (1 gap):             ooooo---ooo
    - transcript:       ...   >>>>>>>>>   >>>>>>>>>    ...
    ```

- "3:jmm:agm:aaf:"

    ```
    - read (2 gaps):             oo---ooooo---o
    - transcript:       ...   >>>>>>>>   >>>>>   >>>>>>>    ...
    ```

For clarity, only the exons involved in the overlap are represented. The transcript can of course have more upstream and downstream exons, which is denoted by the ... on the left side (5' end) and right side (3' end) of each drawing. Note that the exons represented in the 2nd and 3rd drawings are consecutive and adjacent in the processed transcript.

Encodings "1:f:" and "1:j:" are variations of the situation described by encoding "1:i:". For "1:f:", the first aligned base of the read (or "flipped" read) is aligned with the first base of the exon. For "1:j:", the last aligned base of the read (or "flipped" read) is aligned with the last base of the exon:

- "1:f:"

    ```
    - read (no gap):         oooooooo
    - transcript:       ...   >>>>>>>>>>>>>    ...
    ```

- "1:j:"

    ```
    - read (no gap):                oooooooo
    - transcript:       ...   >>>>>>>>>>>>>    ...
    ```

```
> U1.OV00_is_comp <- isCompatibleWithSplicing(U1.ovenc)
> table(U1.OV00_is_comp)  # 531797 "compatible" overlaps

U1.OV00_is_comp
 FALSE    TRUE
 31755  531797
```

Finally, let's extract the "compatible" overlaps from U1.OV00:

```
> U1.compOV00 <- U1.OV00[U1.OV00_is_comp]
```

Note that high-level convenience wrapper findCompatibleOverlaps can be used for computing the "compatible" overlaps directly between a *GAlignments* object (containing reads) and a *GRangesList* object (containing transcripts):

```
> U1.compOV00_again <- findCompatibleOverlaps(U1.GAL, exbytx)
> stopifnot(identical(U1.compOV00_again, U1.compOV00))
```

6.1.2 Tabulate the "compatible" single-end overlaps

Number of "compatible" transcripts for each alignment in U1.GAL:

```
> U1.GAL_ncomptx <- nhitPerQuery(U1.compOV00)
> mcols(U1.GAL)$ncomptx <- U1.GAL_ncomptx
> head(U1.GAL)

GAlignments with 6 alignments and 4 metadata columns:
                  seqnames strand       cigar    qwidth      start       end     width      ngap |
                     <Rle>  <Rle> <character> <integer> <integer> <integer> <integer> <integer> |
  SRR031729.3941844     chr4      -         75M        75       892       966        75         0 |
  SRR031728.3674563     chr4      -         75M        75       919       993        75         0 |
  SRR031729.8532600     chr4      +         75M        75       924       998        75         0 |
  SRR031729.2779333     chr4      +         75M        75       936      1010        75         0 |
  SRR031728.2826481     chr4      +         75M        75       949      1023        75         0 |
  SRR031728.2919098     chr4      -         75M        75       967      1041        75         0 |
                                                                                              seq
                                                                                    <DNAStringSet>
  SRR031729.3941844 CTGTGGTGACCAACACCACAGAATGGTTCGGGCCCAATTAGAGGGTTCCCTGCCCCTTTCCTGGCTAGGTTGTCC
  SRR031728.3674563 TCGGGCCCAATTAGAGGGTTCCCTGCCCCTTTCCTGGCTAGGTTGTCCGCTAGCTCATTTCCTGGGCTGTTGTTG
  SRR031729.8532600 CCCAATTAGAGGATTCTCTGCCCCTTTCCTGGCTAGGTTGTCCGGTAGCTCATTTCCCGGGATGTTGTTGTGTCC
  SRR031729.2779333 GTTCTCTGCCCCTTTCCTGGCTAGGTTGTCCGCTAGCTCATTTCCCGAGATGTTGTTGTGTCCCGGGACCCACCT
  SRR031728.2826481 TTCCTGGCTAGGTTGTCCGCTAGCTCATTTCCCGGGCTGTTGTTGTGTCCCGGGACACACCTTATTGTGAGTTTG
  SRR031728.2919098 GCTAGCTCATTTCCCGGGAGGTTGTTGTGTCCCGGGACCCACCTTATTGTGAGTTTGTTGACAGCTCCAAGTTTG
                           NM       ntx   ncomptx
                    <integer> <integer> <integer>
  SRR031729.3941844         1         0         0
  SRR031728.3674563         3         0         0
  SRR031729.8532600         2         0         0
  SRR031729.2779333         1         0         0
  SRR031728.2826481         2         0         0
  SRR031728.2919098         1         0         0
  ---
  seqlengths:
       chr2L     chr2R     chr3L     chr3R      chr4      chrM      chrX   chrYHet
    23011544  21146708  24543557  27905053   1351857     19517  22422827    347038
```

```
> table(U1.GAL_ncomptx)

U1.GAL_ncomptx
    0     1     2     3     4     5     6     7     8     9    10    11    12
51101  9848 33697 72987  5034 14021  7516   581  1789  2015   530  4389   847
```

```
> mean(U1.GAL_ncomptx >= 1)

[1] 0.7499401
```

75% of the alignments in U1.GAL are "compatible" with at least 1 transcript in exbytx.

Note that high-level convenience wrapper countCompatibleOverlaps can be used directly on U1.GAL and exbytx for computing U1.GAL_ncomptx:

```
> U1.GAL_ncomptx_again <- countCompatibleOverlaps(U1.GAL, exbytx)
> stopifnot(identical(U1.GAL_ncomptx_again, U1.GAL_ncomptx))
```

Number of "compatible" transcripts for each read:

```
> U1.compOV10 <- remapHits(U1.compOV00, query.map=U1.GAL_qnames)
> U1.uqnames_ncomptx <- nhitPerQuery(U1.compOV10)
> names(U1.uqnames_ncomptx) <- U1.uqnames
> table(U1.uqnames_ncomptx)
```

```
U1.uqnames_ncomptx
    0     1     2     3     4     5     6     7     8     9    10    11    12
42886  9711 26075 72989  5413 14044  8584   581  2706  2015   530  4389   847
```

```
> mean(U1.uqnames_ncomptx >= 1)
```

```
[1] 0.7751953
```

77.5% of the reads are "compatible" with at least 1 transcript in exbytx.

Number of "compatible" reads for each transcript:

```
> U1.exbytx_ncompOV10 <- nhitPerSubject(U1.compOV10)
> names(U1.exbytx_ncompOV10) <- names(exbytx)
> mean(U1.exbytx_ncompOV10 >= 50)
```

```
[1] 0.008706681
```

Only 0.87% of the transcripts in exbytx are "compatible" with at least 50 reads.

Top 10 transcripts:

```
> head(sort(U1.exbytx_ncompOV10, decreasing=TRUE), n=10)
```

```
FBtr0308296 FBtr0089175 FBtr0089176 FBtr0089243 FBtr0289951 FBtr0112904 FBtr0089186 FBtr0089187
      40309       40158       33490       11365       11332       11284       10018        9627
FBtr0333672 FBtr0089172
       9568        6599
```

Note that this "top 10" is slightly different from the "top 10" we obtained earlier when we counted **all** the overlaps.

6.2   "Compatible" paired-end overlaps

6.2.1   "Compatible" paired-end encodings

U3.ovenc contains 13 unique paired-end encodings "compatible" with the splicing of the transcript:

```
> sort(U3.ovenc_table[isCompatibleWithSplicing(U3.unique_encodings)])
```

```
      1--2:f--jm:a--af:                   1--1:f--j:               2--1:jm--m:af--j:
                    3                           12                               21
      2--1:jm--m:af--f:        1--1:j--m:a--i: 2--2:jm--mm:af--jm:aa--af:
                   24                           51                              153
      1--1:i--m:a--i:                   1--1:i--j:                     1--1:f--i:
                  287                          403                              617
      1--2:i--jm:a--af:        2--1:jm--m:af--i:                      1--1:i--i:
                 2480                         2700                           100084
```

Paired-end encodings "1--1:i--i:" (100084 occurences in U3.ovenc), "2--1:jm--m:af--i:" (2700 occurences in U3.ovenc), "1--2:i--jm:a--af:" (2480 occurences in U3.ovenc), "1--1:i--m:a--i:" (287 occurences in U3.ovenc), and "2--2:jm--mm:af--jm:aa--af:" (153 occurences in U3.ovenc), correspond to the following paired-end overlaps:

- "1--1:i--i:"

  - paired-end read (no gap on the first end, no gap on the
    last end):   oooo   oooo
  - transcript: ...   >>>>>>>>>>>>>>>>   ...

- "2--1:jm--m:af--i:"

  - paired-end read (1 gap on the first end, no gap on the
    last end):   ooo---o   oooo
  - transcript: ...   >>>>>>>>   >>>>>>>>>>>   ...

- "1--2:i--jm:a--af:"

  - paired-end read (no gap on the first end, 1 gap on the
    last end):   oooo   oo---oo
  - transcript: ... >>>>>>>>>>>>>>   >>>>>>>>>   ...

- "1--1:i--m:a--i:"

  - paired-end read (no gap on the first end, no gap on the
    last end):   oooo   oooo
  - transcript: ...   >>>>>>>>>   >>>>>>>   ...

- "2--2:jm--mm:af--jm:aa--af:"

  - paired-end read (1 gap on the first end, 1 gap on the
    last end):   ooo---o   oo---oo
  - transcript: ...   >>>>>>   >>>>>>>   >>>>>   ...

Note: switch use of "first" and "last" above if the read was "flipped".

```
> U3.OV00_is_comp <- isCompatibleWithSplicing(U3.ovenc)
> table(U3.OV00_is_comp)  # 106835 "compatible" paired-end overlaps

U3.OV00_is_comp
 FALSE    TRUE
  6992 106835
```

Finally, let's extract the "compatible" paired-end overlaps from U3.OV00:

```
> U3.compOV00 <- U3.OV00[U3.OV00_is_comp]
```

Note that, like with our single-end reads, high-level convenience wrapper findCompatibleOverlaps can be used for computing the "compatible" paired-end overlaps directly between a *GAlignmentPairs* object (containing paired-end reads) and a *GRangesList* object (containing transcripts):

```
> U3.compOV00_again <- findCompatibleOverlaps(U3.GALP, exbytx)
> stopifnot(identical(U3.compOV00_again, U3.compOV00))
```

6.2.2   Tabulate the "compatible" paired-end overlaps

Number of "compatible" transcripts for each alignment pair in U3.GALP:

```
> U3.GALP_ncomptx <- nhitPerQuery(U3.compOV00)
> mcols(U3.GALP)$ncomptx <- U3.GALP_ncomptx
> head(U3.GALP)
```

```
GAlignmentPairs with 6 alignment pairs and 2 metadata columns:
                  seqnames strand :      ranges --       ranges |     ntx   ncomptx
                     <Rle>  <Rle> :   <IRanges> --    <IRanges> | <integer> <integer>
SRR031715.1138209     chr4      + : [ 169,  205] -- [ 326,  362] |       0        0
 SRR031714.756385     chr4      + : [ 943,  979] -- [1086, 1122] |       0        0
SRR031714.5054563     chr4      + : [ 946,  982] -- [ 986, 1022] |       0        0
SRR031715.1722593     chr4      + : [ 966, 1002] -- [1108, 1144] |       0        0
SRR031715.2202469     chr4      + : [ 966, 1002] -- [1114, 1150] |       0        0
SRR031714.3544437     chr4      - : [1087, 1123] -- [ 963,  999] |       0        0
---
seqlengths:
    chr2L     chr2R     chr3L     chr3R      chr4      chrM      chrX   chrYHet
 23011544  21146708  24543557  27905053   1351857     19517  22422827    347038

> table(U3.GALP_ncomptx)

U3.GALP_ncomptx
    0     1     2     3     4     5     6     7     8     9    10    11    12
13898  2028  8091 14334  1099  2950  1865    84   296   330    88   699    66

> mean(U3.GALP_ncomptx >= 1)

[1] 0.6967356
```

69.7% of the alignment pairs in U3.GALP are "compatible" with at least 1 transcript in exbytx.
Note that high-level convenience wrapper countCompatibleOverlaps can be used directly on U3.GALP and exbytx for computing U3.GALP_ncomptx:

```
> U3.GALP_ncomptx_again <- countCompatibleOverlaps(U3.GALP, exbytx)
> stopifnot(identical(U3.GALP_ncomptx_again, U3.GALP_ncomptx))
```

Number of "compatible" transcripts for each template:

```
> U3.compOV10 <- remapHits(U3.compOV00, query.map=U3.GALP_qnames)
> U3.uqnames_ncomptx <- nhitPerQuery(U3.compOV10)
> names(U3.uqnames_ncomptx) <- U3.uqnames
> table(U3.uqnames_ncomptx)

U3.uqnames_ncomptx
    0     1     2     3     4     5     6     7     8     9    10    11    12
12783  2026  6531 14334  1210  2950  2114    84   444   330    88   699    66

> mean(U3.uqnames_ncomptx >= 1)

[1] 0.7072081
```

70.7% of the templates are "compatible" with at least 1 transcript in exbytx.
Number of "compatible" templates for each transcript:

```
> U3.exbytx_ncompOV10 <- nhitPerSubject(U3.compOV10)
> names(U3.exbytx_ncompOV10) <- names(exbytx)
> mean(U3.exbytx_ncompOV10 >= 50)

[1] 0.007061324
```

Only 0.7% of the transcripts in exbytx are "compatible" with at least 50 templates.
Top 10 transcripts:

```
> head(sort(U3.exbytx_ncompOV10, decreasing=TRUE), n=10)
```

```
FBtr0308296 FBtr0089175 FBtr0089176 FBtr0289951 FBtr0089243 FBtr0112904 FBtr0089187 FBtr0089186
       7425        7419        5227        2686        2684        2640        2257        2250
FBtr0333672 FBtr0310542
       2206        1649
```

Note that this "top 10" is slightly different from the "top 10" we obtained earlier when we counted **all** the paired-end overlaps.

7   Project the alignments on the transcriptome

7.1   Project the single-end alignments on the transcriptome

The extractQueryStartInTranscript function computes for each overlap the position of the *query start* in the transcript:

```
> U1.OV00_qstart <- extractQueryStartInTranscript(U1.grl, exbytx,
+                                           hits=U1.OV00, ovenc=U1.ovenc)
> head(subset(U1.OV00_qstart, U1.OV00_is_comp))
```

```
   startInTranscript firstSpannedExonRank startInFirstSpannedExon
1               100                    1                      100
8              4229                    5                      137
9              4229                    5                      137
10             4207                    5                      115
11             4207                    5                      115
12             4187                    5                       95
```

U1.OV00_qstart is a data frame with 1 row per overlap and 3 columns:
  1. startInTranscript: the 1-based start position of the read with respect to the transcript. Position 1 always corresponds to the first base on the 5' end of the transcript sequence.
  2. firstSpannedExonRank: the rank of the first exon spanned by the read, that is, the rank of the exon found at position startInTranscript in the transcript.
  3. startInFirstSpannedExon: the 1-based start position of the read with respect to the first exon spanned by the read.

Having this information allows us for example to compare the read and transcript nucleotide sequences for each "compatible" overlap. If we use the *reference query sequence* instead of the *original query sequence* for this comparison, then it should match **exactly** the sequence found at the *query start* in the transcript.

Let's start by using the extractTranscriptsFromGenome to extract the transcript sequences (aka transcriptome) from the dm3 reference genome:

```
> txseq <- extractTranscriptsFromGenome(Dmelanogaster, exbytx)
```

For each "compatible" overlap, the read sequence in U1.GAL_rqseq must be an *exact* substring of the transcript sequence in exbytx_seq:

```
> U1.OV00_rqseq <- U1.GAL_rqseq[queryHits(U1.OV00)]
> U1.OV00_rqseq[flippedQuery(U1.ovenc)] <- reverseComplement(U1.OV00_rqseq[flippedQuery(U1.ovenc)])
> U1.OV00_txseq <- txseq[subjectHits(U1.OV00)]
> stopifnot(all(
+     U1.OV00_rqseq[U1.OV00_is_comp] ==
+         narrow(U1.OV00_txseq[U1.OV00_is_comp],
+             start=U1.OV00_qstart$startInTranscript[U1.OV00_is_comp],
+             width=width(U1.OV00_rqseq)[U1.OV00_is_comp])
+ ))
```

Because of this relationship between the *reference query sequence* and the transcript sequence of a "compatible" overlap, and because of the relationship between the *original query sequences* and the *reference query sequences*, then the edit distance reported in the NM tag is actually the edit distance between the *original query* and the transcript of a "compatible" overlap.

7.2   Project the paired-end alignments on the transcriptome

For a paired-end read, the *query start* is the start of its "left end".

```
> U3.OV00_Lqstart <- extractQueryStartInTranscript(U3.grl, exbytx,
+                                                  hits=U3.OV00, ovenc=U3.ovenc)
> head(subset(U3.OV00_Lqstart, U3.OV00_is_comp))
```

	startInTranscript	firstSpannedExonRank	startInFirstSpannedExon
2	4118	5	26
7	3940	4	31
8	3940	4	31
9	3692	3	320
10	3692	3	320
11	3690	3	318

Note that `extractQueryStartInTranscript` can be called with `for.query.right.end=TRUE` if we want this information for the "right ends" of the reads:

```
> U3.OV00_Rqstart <- extractQueryStartInTranscript(U3.grl, exbytx,
+                                                  hits=U3.OV00, ovenc=U3.ovenc,
+                                                  for.query.right.end=TRUE)
> head(subset(U3.OV00_Rqstart, U3.OV00_is_comp))
```

	startInTranscript	firstSpannedExonRank	startInFirstSpannedExon
2	4267	5	175
7	3948	4	39
8	3948	4	39
9	3849	3	477
10	3849	3	477
11	3831	3	459

Like with single-end reads, having this information allows us for example to compare the read and transcript nucleotide sequences for each "compatible" overlap. If we use the *reference query sequence* instead of the *original query sequence* for this comparison, then it should match **exactly** the sequences of the "left" and "right" ends of the read in the transcript.
Let's assign the "left and right reference query sequences" to each overlap:

```
> U3.OV00_Lrqseq <- U3.GALP_rqseq1[queryHits(U3.OV00)]
> U3.OV00_Rrqseq <- U3.GALP_rqseq2[queryHits(U3.OV00)]
```

For the single-end reads, the sequence associated with a "flipped query" just needed to be "reverse complemented". For paired-end reads, we also need to swap the 2 sequences in the pair:

```
> flip_idx <- which(flippedQuery(U3.ovenc))
> tmp <- U3.OV00_Lrqseq[flip_idx]
> U3.OV00_Lrqseq[flip_idx] <- reverseComplement(U3.OV00_Rrqseq[flip_idx])
> U3.OV00_Rrqseq[flip_idx] <- reverseComplement(tmp)
```

Let's assign the transcript sequence to each overlap:

```
> U3.OV00_txseq <- txseq[subjectHits(U3.OV00)]
```

For each "compatible" overlap, we expect the "left and right reference query sequences" of the read to be *exact* substrings of the transcript sequence. Let's check the "left reference query sequences":

```
> stopifnot(all(
+     U3.OV00_Lrqseq[U3.OV00_is_comp] ==
+         narrow(U3.OV00_txseq[U3.OV00_is_comp],
+                 start=U3.OV00_Lqstart$startInTranscript[U3.OV00_is_comp],
+                 width=width(U3.OV00_Lrqseq)[U3.OV00_is_comp])
+ ))
```

and the "right reference query sequences":

```
> stopifnot(all(
+     U3.OV00_Rrqseq[U3.OV00_is_comp] ==
+         narrow(U3.OV00_txseq[U3.OV00_is_comp],
+                 start=U3.OV00_Rqstart$startInTranscript[U3.OV00_is_comp],
+                 width=width(U3.OV00_Rrqseq)[U3.OV00_is_comp])
+ ))
```

8   Align the reads to the transcriptome

Aligning the reads to the reference genome is not the most efficient nor accurate way to count the number of "compatible" overlaps per *original query*. Supporting junction reads (i.e. reads that align with at least 1 gap) introduces a significant computational cost during the alignment process. Then, as we've seen in the previous sections, each alignment produced by the aligner needs to be broken into a set of ranges (based on its CIGAR) and those ranges compared to the ranges of the exons grouped by transcript.

A more straightforward and accurate approach is to align the reads directly to the transcriptome, and without allowing the typical gap that the aligner needs to introduce when aligning a junction read to the reference genome. With this approach, a "hit" between a read and a transcript is necessarily compatible with the splicing of the transcript. In case of a "hit", we'll say that the read and the transcript are "string-based compatible" (to differentiate from our previous notion of "compatible" overlaps that we will call "encoding-based compatible" from now on, unless the context is clear).

8.1   Align the single-end reads to the transcriptome

8.1.1   Find the "hits"

The single-end reads are in U1.oqseq, the transcriptome is in exbytx_seq.

Since indels were not allowed/supported during the alignment of the reads to the reference genome, we don't need to allow/support them either for aligning the reads to the transcriptome. Also since our goal is to find (and count) "compatible" overlaps between reads and transcripts, we don't need to keep track of the details of the alignments between the reads and the transcripts. Finally, since BAM file untreated1_chr4.bam is not the full output of the aligner but the subset obtained by keeping only the alignments located on chr4, we don't need to align U1.oqseq to the full transcriptome, but only to the subset of exbytx_seq made of the transcripts located on chr4.

With those simplifications in mind, we write the following function that we will use to find the "hits" between the reads and the transcriptome:

```
> ### A wrapper to vwhichPDict() that supports IUPAC ambiguity codes in 'qseq'
> ### and 'txseq', and treats them as such.
> findSequenceHits <- function(qseq, txseq, which.txseq=NULL, max.mismatch=0)
+ {
+     .asHits <- function(x, pattern_length)
+     {
+         query_hits <- unlist(x)
+         if (is.null(query_hits))
+             query_hits <- integer(0)
```

```
+          subject_hits <- rep.int(seq_len(length(x)), elementLengths(x))
+          new("Hits", queryHits=query_hits, subjectHits=subject_hits,
+                       queryLength=pattern_length, subjectLength=length(x))
+      }
+
+      .isHitInTranscriptBounds <- function(hits, qseq, txseq)
+      {
+          sapply(seq_len(length(hits)),
+              function(i) {
+                  pattern <- qseq[[queryHits(hits)[i]]]
+                  subject <- txseq[[subjectHits(hits)[i]]]
+                  v <- matchPattern(pattern, subject,
+                                       max.mismatch=max.mismatch, fixed=FALSE)
+                  any(1L <= start(v) & end(v) <= length(subject))
+              })
+      }
+
+      if (!is.null(which.txseq)) {
+          txseq0 <- txseq
+          txseq <- txseq[which.txseq]
+      }
+
+      names(qseq) <- NULL
+      other <- alphabetFrequency(qseq, baseOnly=TRUE)[ , "other"]
+      is_clean <- other == 0L  # "clean" means "no IUPAC ambiguity code"
+
+      ## Find hits for "clean" original queries.
+      qseq0 <- qseq[is_clean]
+      pdict0 <- PDict(qseq0, max.mismatch=max.mismatch)
+      m0 <- vwhichPDict(pdict0, txseq,
+                      max.mismatch=max.mismatch, fixed="pattern")
+      hits0 <- .asHits(m0, length(qseq0))
+      hits0@queryLength <- length(qseq)
+      hits0@queryHits <- which(is_clean)[hits0@queryHits]
+
+      ## Find hits for non "clean" original queries.
+      qseq1 <- qseq[!is_clean]
+      m1 <- vwhichPDict(qseq1, txseq,
+                      max.mismatch=max.mismatch, fixed=FALSE)
+      hits1 <- .asHits(m1, length(qseq1))
+      hits1@queryLength <- length(qseq)
+      hits1@queryHits <- which(!is_clean)[hits1@queryHits]
+
+      ## Combine the hits.
+      query_hits <- c(queryHits(hits0), queryHits(hits1))
+      subject_hits <- c(subjectHits(hits0), subjectHits(hits1))
+
+      if (!is.null(which.txseq)) {
+          ## Remap the hits.
+          txseq <- txseq0
+          subject_hits <- which.txseq[subject_hits]
+          hits0@subjectLength <- length(txseq)
+      }
+
```

```
+       ## Order the hits.
+       oo <- IRanges:::orderIntegerPairs(query_hits, subject_hits)
+       hits0@queryHits <- query_hits[oo]
+       hits0@subjectHits <- subject_hits[oo]
+
+       if (max.mismatch != 0L) {
+           ## Keep only "in bounds" hits.
+           is_in_bounds <- .isHitInTranscriptBounds(hits0, qseq, txseq)
+           hits0 <- hits0[is_in_bounds]
+       }
+       hits0
+ }
```

Let's compute the index of the transcripts in exbytx_seq located on chr4 (findSequenceHits will restrict the search to those transcripts):

```
> chr4tx <- transcripts(txdb, vals=list(tx_chrom="chr4"))
> chr4txnames <- mcols(chr4tx)$tx_name
> which.txseq <- match(chr4txnames, names(txseq))
```

We know that the aligner tolerated up to 6 mismatches per read. The 3 following commands find the "hits" for each *original query*, then find the "hits" for each "flipped *original query*", and finally merge all the "hits" (note that the 3 commands take about 1 hour to complete on a modern laptop):

```
> U1.sbcompHITSa <- findSequenceHits(U1.oqseq, txseq,
+                                    which.txseq=which.txseq, max.mismatch=6)
> U1.sbcompHITSb <- findSequenceHits(reverseComplement(U1.oqseq), txseq,
+                                    which.txseq=which.txseq, max.mismatch=6)
> U1.sbcompHITS <- union(U1.sbcompHITSa, U1.sbcompHITSb)
```

8.1.2 Tabulate the "hits"

Number of "string-based compatible" transcripts for each read:

```
> U1.uqnames_nsbcomptx <- nhitPerQuery(U1.sbcompHITS)
> names(U1.uqnames_nsbcomptx) <- U1.uqnames
> table(U1.uqnames_nsbcomptx)

U1.uqnames_nsbcomptx
    0     1     2     3     4     5     6     7     8     9    10    11    12
40555 10080 25299 74609  5207 14265  8643   610  3410  2056   534  4588   914

> mean(U1.uqnames_nsbcomptx >= 1)

[1] 0.7874142
```

77.7% of the reads are "string-based compatible" with at least 1 transcript in exbytx.
Number of "string-based compatible" reads for each transcript:

```
> U1.exbytx_nsbcompHITS <- nhitPerSubject(U1.sbcompHITS)
> names(U1.exbytx_nsbcompHITS) <- names(exbytx)
> mean(U1.exbytx_nsbcompHITS >= 50)

[1] 0.008809516
```

Only 0.865% of the transcripts in exbytx are "string-based compatible" with at least 50 reads.
Top 10 transcripts:

```
> head(sort(U1.exbytx_nsbcompHITS, decreasing=TRUE), n=10)
```

```
FBtr0308296 FBtr0089175 FBtr0089176 FBtr0089243 FBtr0289951 FBtr0112904 FBtr0089186 FBtr0333672
      40548       40389       34275       11605       11579       11548       10059        9742
FBtr0089187 FBtr0089172
       9666        6704
```

8.1.3   A closer look at the "hits"

[WORK IN PROGRESS, might be removed or replaced soon...]
    Any "encoding-based compatible" overlap is of course "string-based compatible":

```
> stopifnot(length(setdiff(U1.compOV10, U1.sbcompHITS)) == 0)
```

but the reverse is not true:

```
> length(setdiff(U1.sbcompHITS, U1.compOV10))
```

```
[1] 13549
```

8.2   Align the paired-end reads to the transcriptome

[COMING SOON...]

9   "Almost compatible" overlaps

In many aspects, "compatible" overlaps can be seen as perfect. We are now insterested in a less perfect type of
overlap where the read overlaps the transcript in a way that *would* be "compatible" if 1 or more exons were re-
moved from the transcript. In that case we say that the overlap is "almost compatible" with the transcript. The
isCompatibleWithSkippedExons function can be used on an *OverlapEncodings* object to detect this type of overlap.
Note that isCompatibleWithSkippedExons can also be used on a character vector of factor.

9.1   "Almost compatible" single-end overlaps

9.1.1   "Almost compatible" single-end encodings

U1.ovenc contains 7 unique encodings "almost compatible" with the splicing of the transcript:

```
> sort(U1.ovenc_table[isCompatibleWithSkippedExons(U1.unique_encodings)])
```

```
   2:jm:am:am:am:am:af: 2:jm:am:am:am:am:am:af:                      2:gm:am:af:        2:jm:am:am:am:af:
                      1                       1                                4                        7
   3:jmm:agm:aam:aam:aaf:       3:jmm:agm:aam:aaf:                   2:jm:am:am:af:             2:jm:am:af:
                      9                      21                              144                     1015
```

    Encodings "2:jm:am:af:" (1015 occurences in U1.ovenc), "2:jm:am:am:af:" (144 occurences in U1.ovenc), and
"3:jmm:agm:aam:aaf:" (21 occurences in U1.ovenc), correspond to the following overlaps:

- "2:jm:am:af:"

```
    - read (1 gap):          ooooo---------ooo
    - transcript:      ...  >>>>>>>   >>>>   >>>>>>>>   ...
```

- "2:jm:am:am:af:"

```
    - read (1 gap):          ooooo----------------ooo
    - transcript:      ...  >>>>>>>   >>>>   >>>>>   >>>>>>>>   ...
```

- "3:jmm:agm:aam:aaf:"

```
        - read (2 gaps):                    oo---oooo-----------oo
        - transcript:        ...  >>>>>>>   >>>>    >>>>>    >>>>>>>>    ...

> U1.OV00_is_acomp <- isCompatibleWithSkippedExons(U1.ovenc)
> table(U1.OV00_is_acomp)  # 1202 "almost compatible" overlaps

U1.OV00_is_acomp
 FALSE    TRUE
562350    1202
```

Finally, let's extract the "almost compatible" overlaps from U1.OV00:

```
> U1.acompOV00 <- U1.OV00[U1.OV00_is_acomp]
```

9.1.2   Tabulate the "almost compatible" single-end overlaps

Number of "almost compatible" transcripts for each alignment in U1.GAL:

```
> U1.GAL_nacomptx <- nhitPerQuery(U1.acompOV00)
> mcols(U1.GAL)$nacomptx <- U1.GAL_nacomptx
> head(U1.GAL)

GAlignments with 6 alignments and 5 metadata columns:
                   seqnames strand        cigar     qwidth      start        end      width       ngap |
                      <Rle>  <Rle> <character>  <integer>  <integer>  <integer>  <integer>  <integer> |
   SRR031729.3941844    chr4      -          75M         75        892        966         75          0 |
   SRR031728.3674563    chr4      -          75M         75        919        993         75          0 |
   SRR031729.8532600    chr4      +          75M         75        924        998         75          0 |
   SRR031729.2779333    chr4      +          75M         75        936       1010         75          0 |
   SRR031728.2826481    chr4      +          75M         75        949       1023         75          0 |
   SRR031728.2919098    chr4      -          75M         75        967       1041         75          0 |
                                                                                                    seq
                                                                                          <DNAStringSet>
   SRR031729.3941844 CTGTGGTGACCAACACCACAGAATGGTTCGGGCCCAATTAGAGGGTTCCCTGCCCCTTTCCTGGCTAGGTTGTCC
   SRR031728.3674563 TCGGGCCCAATTAGAGGGTTCCCTGCCCCTTTCCTGGCTAGGTTGTCCGCTAGCTCATTTCCTGGGCTGTTGTTG
   SRR031729.8532600 CCCAATTAGAGGATTCTCTGCCCCTTTCCTGGCTAGGTTGTCCGGTAGCTCATTTCCCGGGATGTTGTTGTGTCC
   SRR031729.2779333 GTTCTCTGCCCCTTTCCTGGCTAGGTTGTCCGCTAGCTCATTTCCCGAGATGTTGTTGTGTCCCGGGACCCACCT
   SRR031728.2826481 TTCCTGGCTAGGTTGTCCGCTAGCTCATTTCCCGGGCTGTTGTTGTGTCCCGGGACACACCTTATTGTGAGTTTG
   SRR031728.2919098 GCTAGCTCATTTCCCGGGAGGTTGTTGTGTCCCGGGACCCACCTTATTGTGAGTTTGTTGACAGCTCCAAGTTTG
                           NM        ntx    ncomptx   nacomptx
                    <integer>  <integer>  <integer>  <integer>
   SRR031729.3941844        1          0          0          0
   SRR031728.3674563        3          0          0          0
   SRR031729.8532600        2          0          0          0
   SRR031729.2779333        1          0          0          0
   SRR031728.2826481        2          0          0          0
   SRR031728.2919098        1          0          0          0
   ---
   seqlengths:
       chr2L      chr2R      chr3L      chr3R       chr4       chrM       chrX    chrYHet
    23011544   21146708   24543557   27905053    1351857      19517   22422827     347038

> table(U1.GAL_nacomptx)

U1.GAL_nacomptx
     0       1       2       3       4       5       6       7       8       9      10      11      12
203800     283     101     107      19      24       2       3       1       3       4       4       4
```

```
> mean(U1.GAL_nacomptx >= 1)
```

```
[1] 0.002715862
```

Only 0.27% of the alignments in `U1.GAL` are "almost compatible" with at least 1 transcript in `exbytx`.
Number of "almost compatible" alignments for each transcript:

```
> U1.exbytx_nacompOV00 <- nhitPerSubject(U1.acompOV00)
> names(U1.exbytx_nacompOV00) <- names(exbytx)
> table(U1.exbytx_nacompOV00)
```

```
U1.exbytx_nacompOV00
    0     1     2     3     4     5     6     7     8     9    10    12    13    14    17    18
29039    50     8    15    12     2     3     7     5     7     3     2     1     1     1     2
   20    21    32    34    44    55    59    77   170
    1     3     2     1     3     2     1     1     1
```

```
> mean(U1.exbytx_nacompOV00 >= 50)
```

```
[1] 0.0001713914
```

Only 0.017% of the transcripts in `exbytx` are "almost compatible" with at least 50 alignments in `U1.GAL`.
Finally note that the "query start in transcript" values returned by `extractQueryStartInTranscript` are also defined
for "almost compatible" overlaps:

```
> head(subset(U1.OV00_qstart, U1.OV00_is_acomp))
```

```
       startInTranscript firstSpannedExonRank startInFirstSpannedExon
144233               133                    1                     133
144234               133                    1                     133
144247               151                    1                     151
144248               151                    1                     151
146615               757                    7                      39
146616               689                    8                      39
```

9.2   "Almost compatible" paired-end overlaps

9.2.1   "Almost compatible" paired-end encodings

`U3.ovenc` contains 5 unique paired-end encodings "almost compatible" with the splicing of the transcript:

```
> sort(U3.ovenc_table[isCompatibleWithSkippedExons(U3.unique_encodings)])
```

```
    2--1:jm--m:am--m:am--m:af--i:        1--2:i--jm:a--am:a--am:a--af:
                               1                                    5
2--2:jm--mm:am--mm:af--jm:aa--af:           1--2:i--jm:a--am:a--af:
                               9                                   53
          2--1:jm--m:am--m:af--i:
                              73
```

Paired-end encodings `"2--1:jm--m:am--m:af--i:"` (73 occurences in `U3.ovenc`), `"1--2:i--jm:a--am:a--af:"`
(53 occurences in `U3.ovenc`), and `"2--2:jm--mm:am--mm:af--jm:aa--af:"` (9 occurences in `U3.ovenc`), correspond
to the following paired-end overlaps:

- `"2--1:jm--m:am--m:af--i:"`

  - paired-end read (1 gap on the first end, no gap on the
    last end):   ooo---------o   oooo
  - transcript: ...   >>>>>   >>>>   >>>>>>>>>   ...

- "1--2:i--jm:a--am:a--af:"

    - paired-end read (no gap on the first end, 1 gap on the
      last end):   oooo   oo---------oo
    - transcript: ... >>>>>>>>>>   >>>   >>>>>>   ...

- "2--2:jm--mm:am--mm:af--jm:aa--af:"

    - paired-end read (1 gap on the first end, 1 gap on the
      last end):   o----------ooo   oo---oo
    - transcript: ... >>>>>   >>>>   >>>>>>>>   >>>>>>   ...

Note: switch use of "first" and "last" above if the read was "flipped".

```
> U3.OV00_is_acomp <- isCompatibleWithSkippedExons(U3.ovenc)
> table(U3.OV00_is_acomp)  # 141 "almost compatible" paired-end overlaps

U3.OV00_is_acomp
 FALSE    TRUE
113686     141
```

Finally, let's extract the "almost compatible" paired-end overlaps from U3.OV00:

```
> U3.acompOV00 <- U3.OV00[U3.OV00_is_acomp]
```

9.2.2   Tabulate the "almost compatible" paired-end overlaps

Number of "almost compatible" transcripts for each alignment pair in U3.GALP:

```
> U3.GALP_nacomptx <- nhitPerQuery(U3.acompOV00)
> mcols(U3.GALP)$nacomptx <- U3.GALP_nacomptx
> head(U3.GALP)

GAlignmentPairs with 6 alignment pairs and 3 metadata columns:
                  seqnames strand :       ranges --       ranges |       ntx    ncomptx   nacomptx
                     <Rle>  <Rle> :    <IRanges> --    <IRanges> | <integer> <integer>  <integer>
SRR031715.1138209     chr4      + : [ 169,  205] -- [ 326,  362] |         0         0          0
 SRR031714.756385     chr4      + : [ 943,  979] -- [1086, 1122] |         0         0          0
SRR031714.5054563     chr4      + : [ 946,  982] -- [ 986, 1022] |         0         0          0
SRR031715.1722593     chr4      + : [ 966, 1002] -- [1108, 1144] |         0         0          0
SRR031715.2202469     chr4      + : [ 966, 1002] -- [1114, 1150] |         0         0          0
SRR031714.3544437     chr4      - : [1087, 1123] -- [ 963,  999] |         0         0          0
---
seqlengths:
    chr2L     chr2R     chr3L     chr3R      chr4      chrM      chrX   chrYHet
 23011544  21146708  24543557  27905053   1351857     19517  22422827    347038

> table(U3.GALP_nacomptx)

U3.GALP_nacomptx
    0     1     2     3     4     5    11
45734    74     4    13     1     1     1

> mean(U3.GALP_nacomptx >= 1)

[1] 0.002051148
```

Only 0.2% of the alignment pairs in U3.GALP are "almost compatible" with at least 1 transcript in exbytx.
Number of "almost compatible" alignment pairs for each transcript:

```
> U3.exbytx_nacompOV00 <- nhitPerSubject(U3.acompOV00)
> names(U3.exbytx_nacompOV00) <- names(exbytx)
> table(U3.exbytx_nacompOV00)

U3.exbytx_nacompOV00
    0     1     5     8    12    13    66
29143    22     4     1     1     1     1

> mean(U3.exbytx_nacompOV00 >= 50)

[1] 3.427827e-05
```

Only 0.0034% of the transcripts in exbytx are "almost compatible" with at least 50 alignment pairs in U3.GALP.

Finally note that the "query start in transcript" values returned by extractQueryStartInTranscript are also defined for "almost compatible" paired-end overlaps:

```
> head(subset(U3.OV00_Lqstart, U3.OV00_is_acomp))

      startInTranscript firstSpannedExonRank startInFirstSpannedExon
27617              1549                   12                       45
27629              1562                   12                       58
27641              1562                   12                       58
27692              1567                   12                       63
27814              1549                   12                       45
42870               659                    4                      101

> head(subset(U3.OV00_Rqstart, U3.OV00_is_acomp))

      startInTranscript firstSpannedExonRank startInFirstSpannedExon
27617              2135                   14                      115
27629              2141                   14                      121
27641              2135                   14                      115
27692              2048                   14                       28
27814              2136                   14                      116
42870               866                    6                       19
```

10 Detect novel splice junctions

10.1 By looking at single-end overlaps

An alignment in U1.GAL with "almost compatible" overlaps but no "compatible" overlaps suggests the presence of one or more transcripts that are not in our annotations.

First we extract the index of those alignments (*nsj* here stands for "**n**ovel **s**plice **j**unction"):

```
> U1.GAL_is_nsj <- U1.GAL_nacomptx != 0L & U1.GAL_ncomptx == 0L
> head(which(U1.GAL_is_nsj))

[1] 57972 57974 58321 67251 67266 67267
```

We make this an index into U1.OV00:

```
> U1.OV00_is_nsj <- queryHits(U1.OV00) %in% which(U1.GAL_is_nsj)
```

We intersect with U1.OV00_is_acomp and then subset U1.OV00 to keep only the overlaps that suggest novel splicing:

```
> U1.OV00_is_nsj <- U1.OV00_is_nsj & U1.OV00_is_acomp
> U1.nsjOV00 <- U1.OV00[U1.OV00_is_nsj]
```

For each overlap in `U1.nsjOV00`, we extract the ranks of the skipped exons (we use a list for this as there might be more than 1 skipped exon per overlap):

```
> U1.nsjOV00_skippedex <- extractSkippedExonRanks(U1.ovenc)[U1.OV00_is_nsj]
> names(U1.nsjOV00_skippedex) <- queryHits(U1.nsjOV00)
> table(elementLengths(U1.nsjOV00_skippedex))

  1   2   3   4   5
234 116   7   1   1
```

Finally, we split `U1.nsjOV00_skippedex` by transcript names:

```
> f <- factor(names(exbytx)[subjectHits(U1.nsjOV00)], levels=names(exbytx))
> U1.exbytx_skippedex <- split(U1.nsjOV00_skippedex, f)
```

`U1.exbytx_skippedex` is a named list of named lists of integer vectors. The first level of names (outer names) are transcript names and the second level of names (inner names) are alignment indices into `U1.GAL`:

```
> head(names(U1.exbytx_skippedex))  # transcript names

[1] "FBtr0300689" "FBtr0300690" "FBtr0330654" "FBtr0309810" "FBtr0306539" "FBtr0306536"
```

Transcript FBtr0089124 receives 7 hits. All of them skip exons 9 and 10:

```
> U1.exbytx_skippedex$FBtr0089124

$`104549`
[1]  9 10

$`104550`
[1]  9 10

$`104553`
[1]  9 10

$`104557`
[1]  9 10

$`104560`
[1]  9 10

$`104572`
[1]  9 10

$`104577`
[1]  9 10
```

Transcript FBtr0089147 receives 4 hits. Two of them skip exon 2, one of them skips exons 2 to 6, and one of them skips exon 10:

```
> U1.exbytx_skippedex$FBtr0089147

$`72828`
[1] 10

$`74018`
[1] 2 3 4 5 6
```

```
$`74664`
[1] 2

$`74670`
[1] 2
```

A few words about the interpretation of U1.exbytx_skippedex: Because of how we've conducted this analysis, the aligments reported in U1.exbytx_skippedex are guaranteed to not have any "compatible" overlaps with other known transcripts. All we can say, for example in the case of transcript FBtr0089124, is that the 7 reported hits that skip exons 9 and 10 show evidence of one or more unknown transcripts with a splice junction that corresponds to the gap between exons 8 and 11. But without further analysis, we can't make any assumption about the exons structure of those unknown transcripts. In particular, we cannot assume the existence of an unknown transcript made of the same exons as transcript FBtr0089124 minus exons 9 and 10!

10.2   By looking at paired-end overlaps

[COMING SOON...]

11   sessionInfo()

```
> sessionInfo()

R version 3.0.2 (2013-09-25)
Platform: x86_64-unknown-linux-gnu (64-bit)

locale:
 [1] LC_CTYPE=en_US.UTF-8        LC_NUMERIC=C               LC_TIME=en_US.UTF-8
 [4] LC_COLLATE=C               LC_MONETARY=en_US.UTF-8    LC_MESSAGES=en_US.UTF-8
 [7] LC_PAPER=en_US.UTF-8       LC_NAME=C                  LC_ADDRESS=C
[10] LC_TELEPHONE=C             LC_MEASUREMENT=en_US.UTF-8 LC_IDENTIFICATION=C

attached base packages:
[1] parallel   stats      graphics   grDevices utils      datasets   methods    base

other attached packages:
 [1] BSgenome.Dmelanogaster.UCSC.dm3_1.3.19      BSgenome.Scerevisiae.UCSC.sacCer2_1.3.19
 [3] KEGGgraph_1.20.0                            graph_1.40.0
 [5] XML_3.98-1.1                                KEGG.db_2.10.1
 [7] RSQLite_0.11.4                              DBI_0.2-7
 [9] BSgenome.Hsapiens.UCSC.hg19_1.3.19          BSgenome_1.30.0
[11] TxDb.Hsapiens.UCSC.hg19.knownGene_2.10.1    edgeR_3.4.0
[13] limma_3.18.0                                DESeq_1.14.0
[15] lattice_0.20-24                             locfit_1.5-9.1
[17] AnnotationHub_1.2.0                         TxDb.Dmelanogaster.UCSC.dm3.ensGene_2.10.1
[19] GenomicFeatures_1.14.0                      AnnotationDbi_1.24.0
[21] Biobase_2.22.0                              pasillaBamSubset_0.0.8
[23] Rsamtools_1.14.1                            Biostrings_2.30.0
[25] GenomicRanges_1.14.1                        XVector_0.2.0
[27] IRanges_1.20.0                              BiocGenerics_0.8.0

loaded via a namespace (and not attached):
 [1] BiocInstaller_1.12.0 BiocStyle_1.0.0      RColorBrewer_1.0-5   RCurl_1.95-4.1
 [5] annotate_1.40.0      biomaRt_2.18.0       bitops_1.0-6         genefilter_1.44.0
```

```
 [9] geneplotter_1.40.0   grid_3.0.2      rjson_0.2.13     rtracklayer_1.22.0
[13] splines_3.0.2        stats4_3.0.2    survival_2.37-4  tools_3.0.2
[17] xtable_1.7-1         zlibbioc_1.8.0
```